



# Architecture

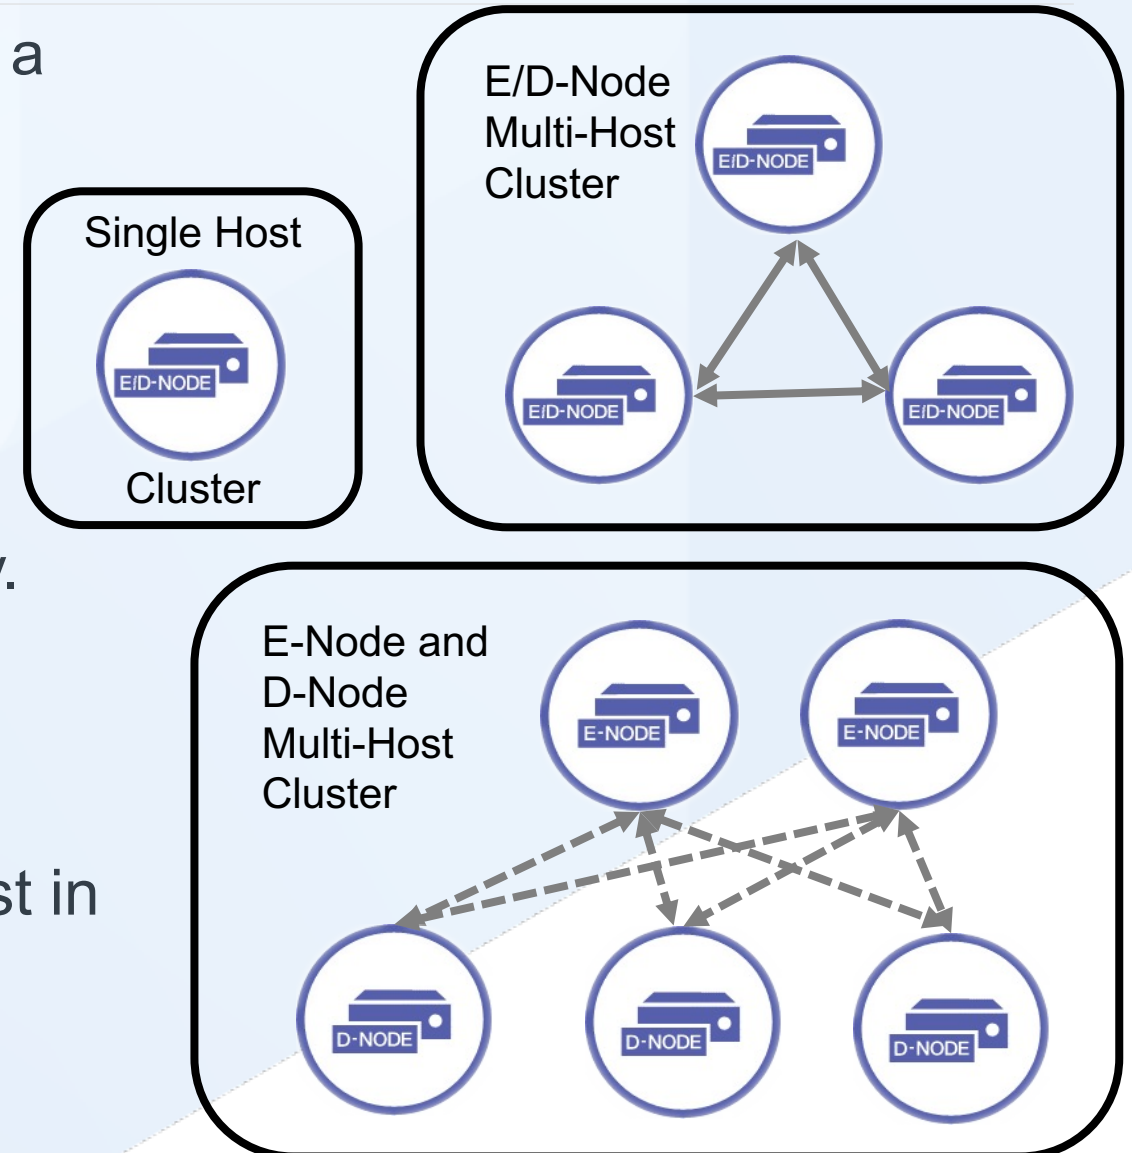
# Key Concepts in this Unit

---

- What a MarkLogic Cluster is.
- How a database distributes documents on Hosts in a Cluster.
- How content is stored.
- Role of the Evaluator Node and Data Node.
- Communicating with a database.

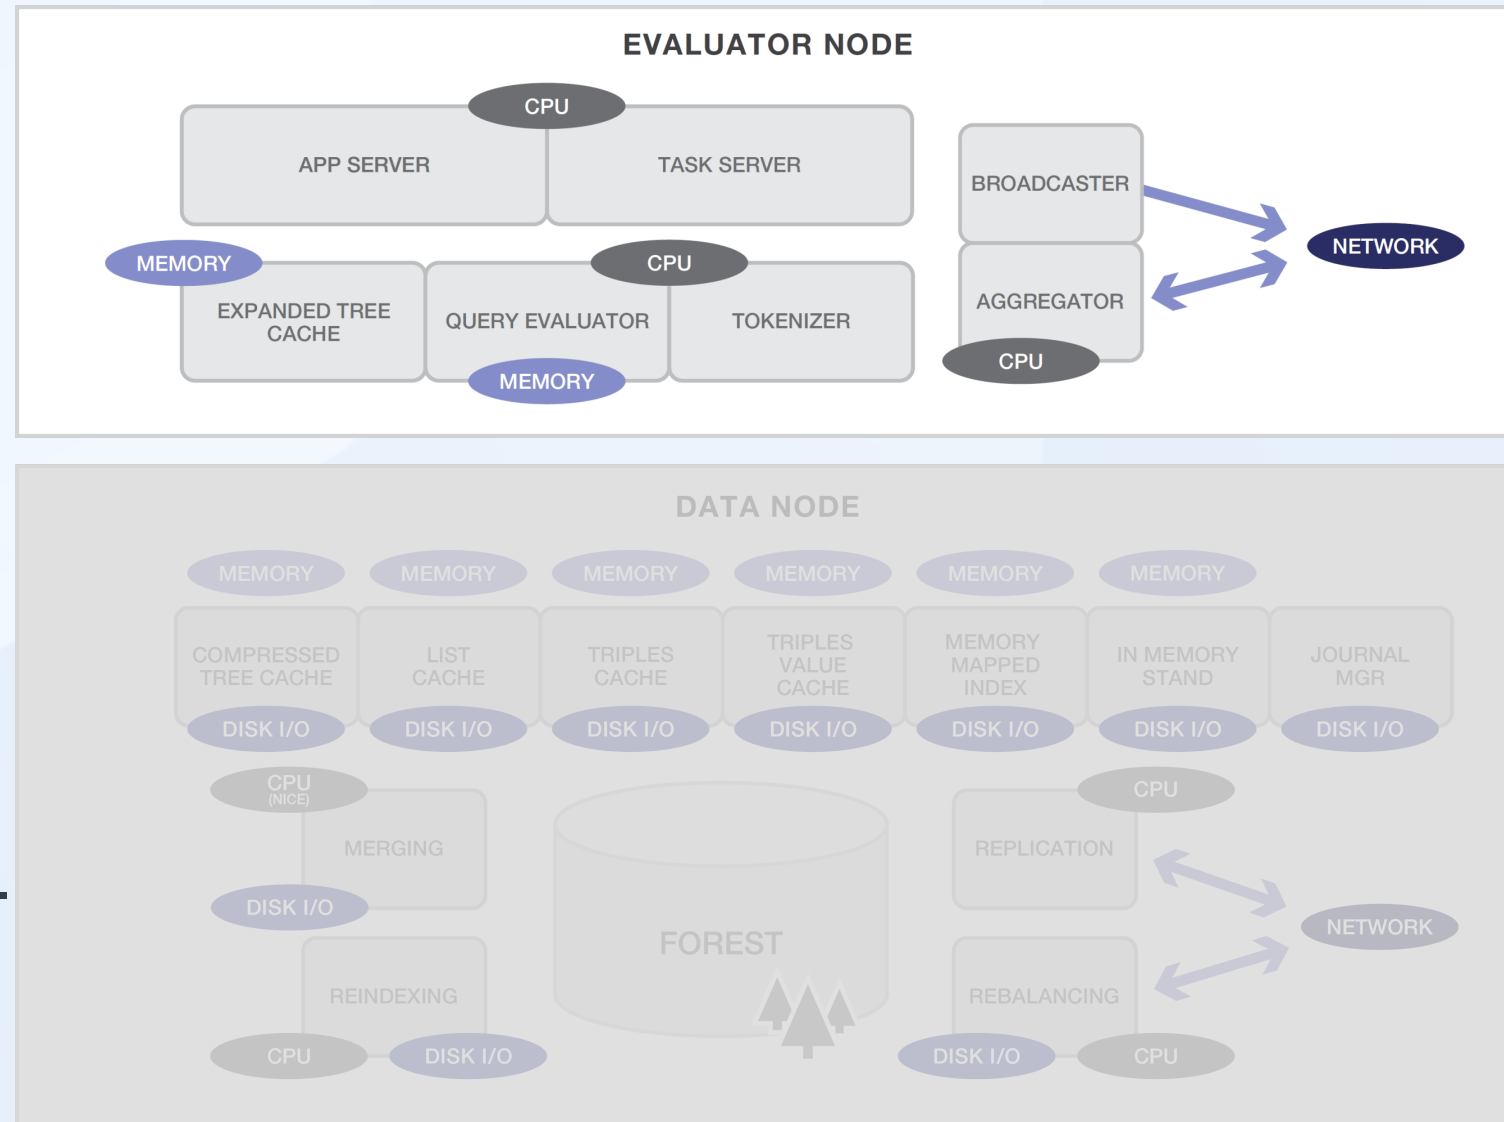
# MarkLogic Cluster

- Consists of 1 or more Hosts/Servers running a MarkLogic Server.
- Each Host can be a/an:
  - Evaluator Node (E-Node).
  - Data Node (D-Node).
  - Both E-Node and D-Node simultaneously.
- Mid-sized Cluster might contain:
  - 2 E-Nodes and 10 D-Nodes.
- Shared-nothing architecture - no single host in charge.



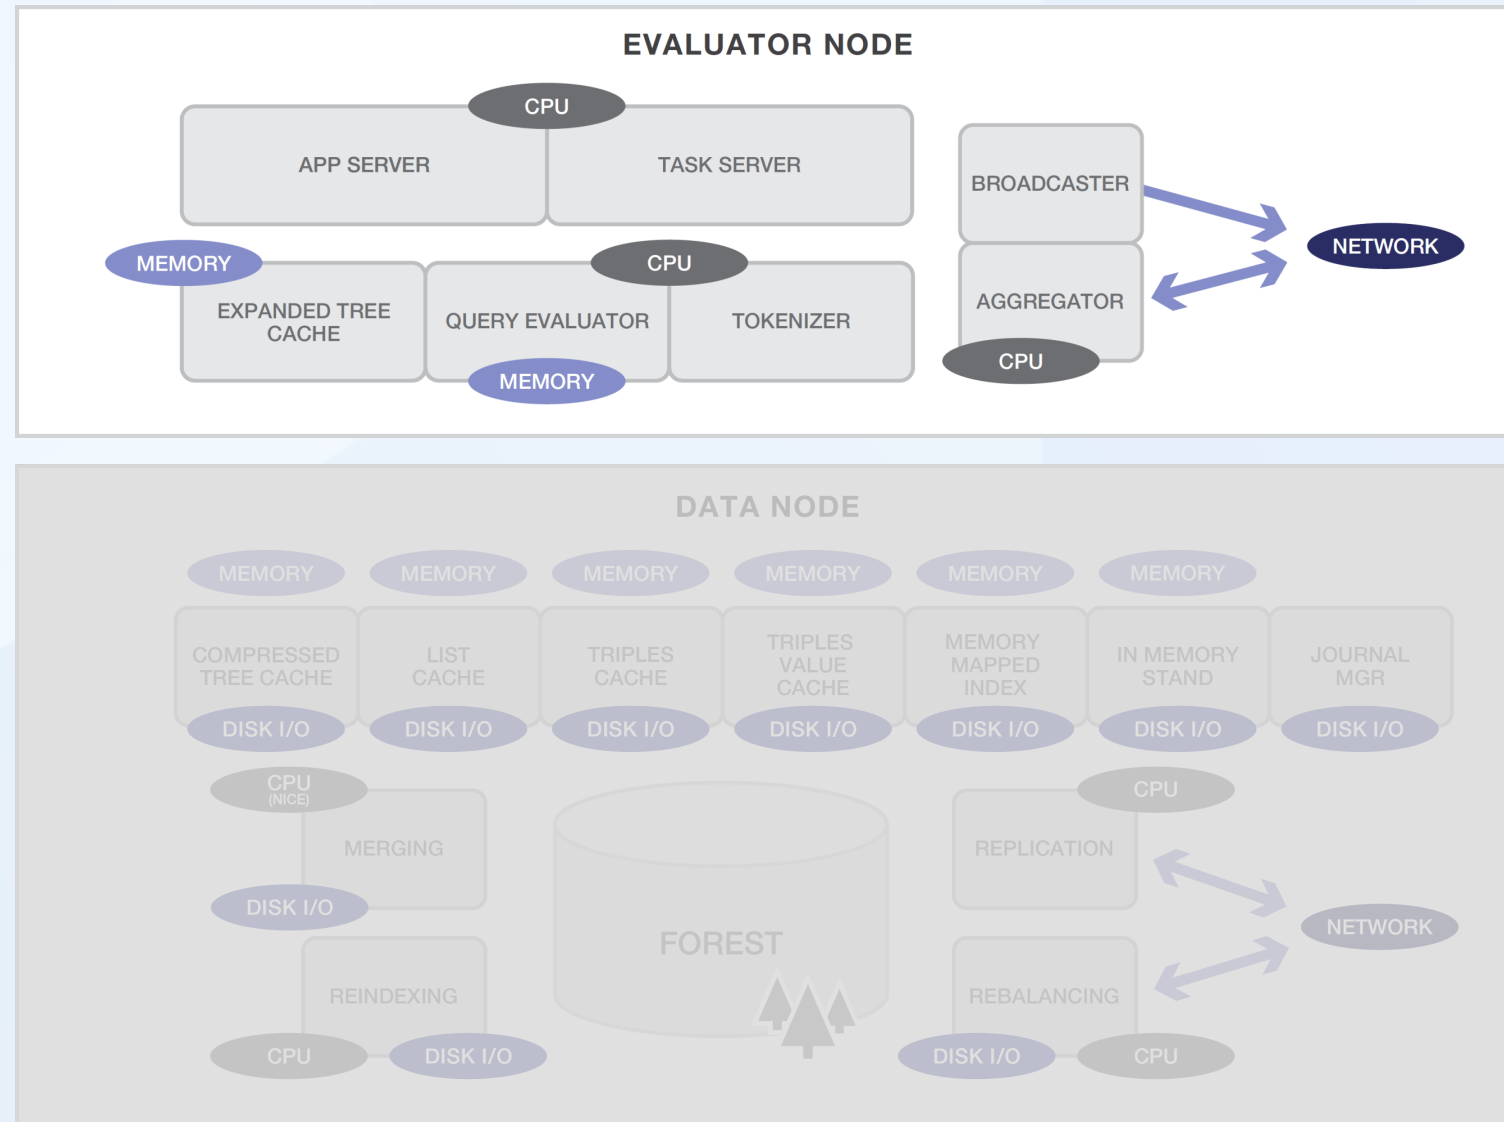
# E-Node Tasks

- App Server listening on a specific port.
- Parses incoming requests/queries.
- Communicates with D-Nodes.
- Generates responses.
- Performs request/response transformations.
- Manages Host memory/cache.



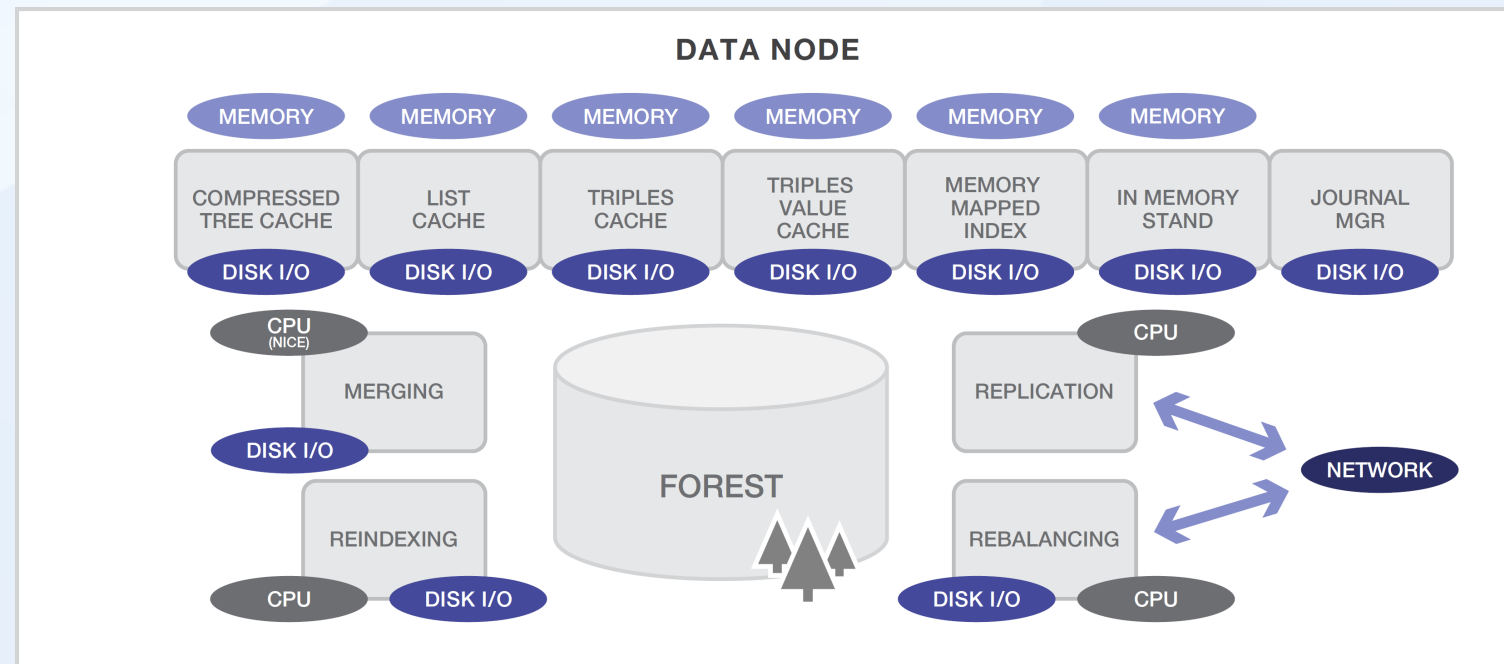
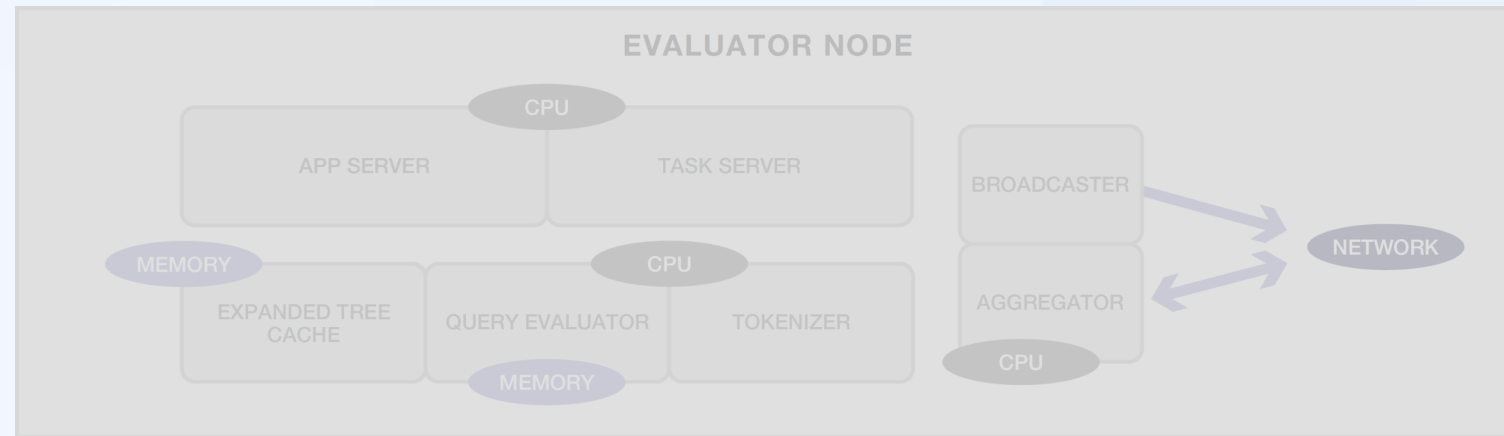
# E-Node Tasks

- Only one E-node is involved per transaction/request.
- The **ONLY** way to communicate with a MarkLogic database.
- Cache settings invest higher memory on “Expanded Tree Cache”.



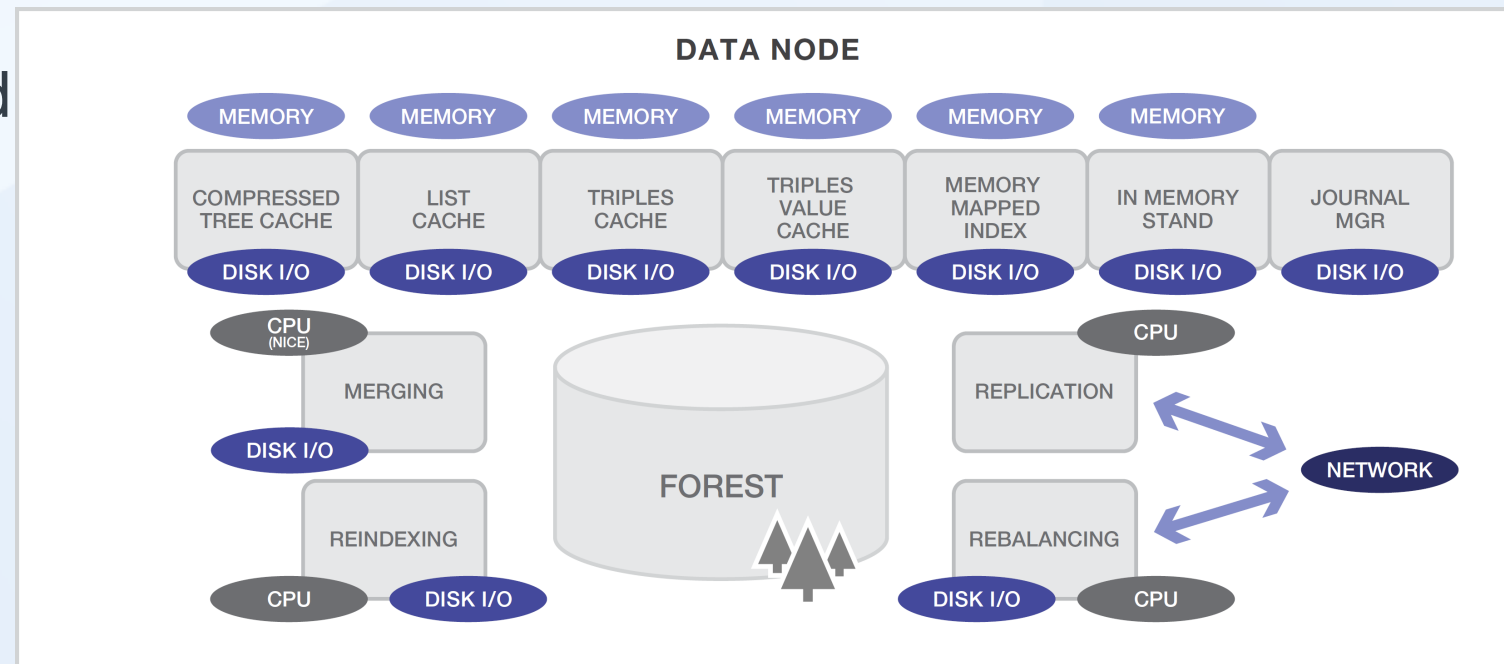
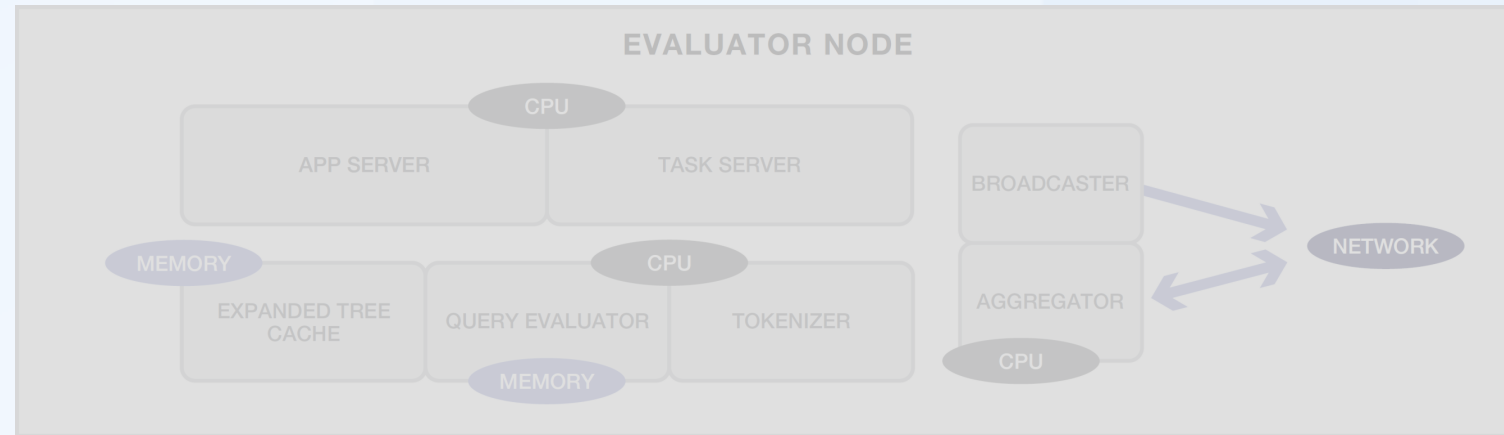
# D-Node Tasks

- Securely stores data in Forests in a compressed, proprietary format.
- Creates and maintains Indexes and Word Lists.
- Provides data to E-Nodes.
- Manages Host memory/cache.
- Performs Merging, Reindexing, Replication and Rebalancing.
- Replies **ONLY** to E-Node requests.



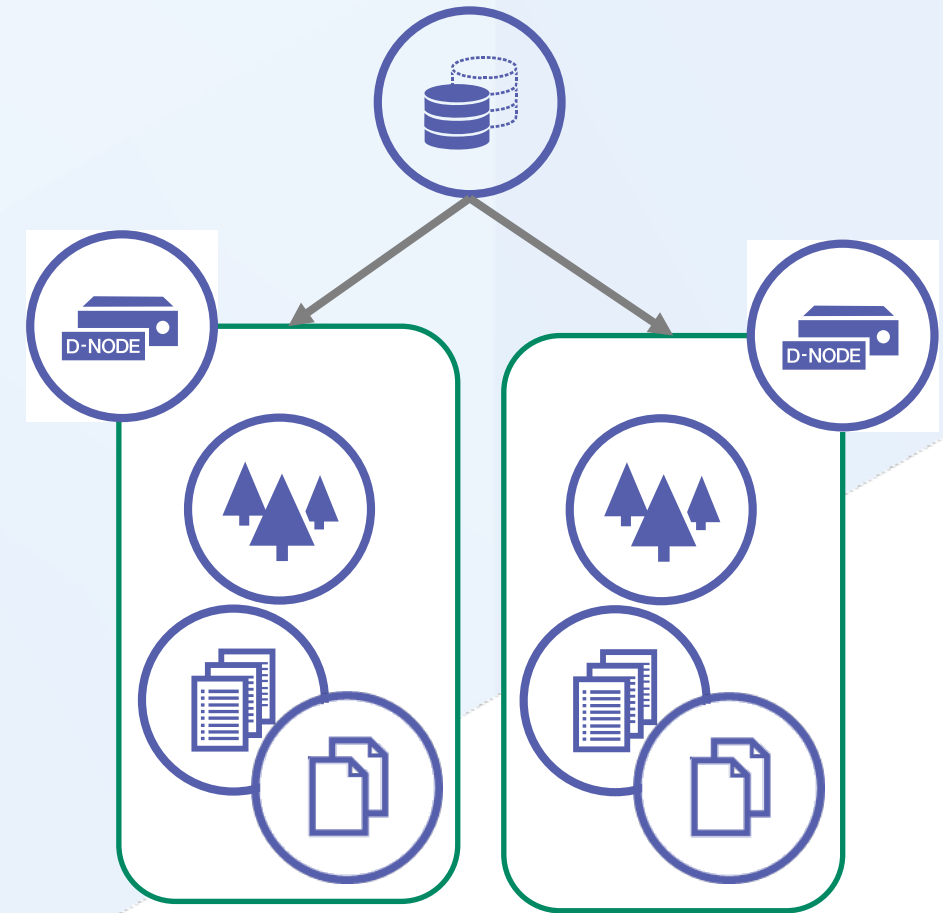
# D-Node Tasks

- All D-Nodes are involved in processing a transaction.
- Cache settings invest higher memory on “List Cache” and “Compressed Tree Cache”.
- More information can be found in [MarkLogic Concepts Guide](#).



# MarkLogic Database

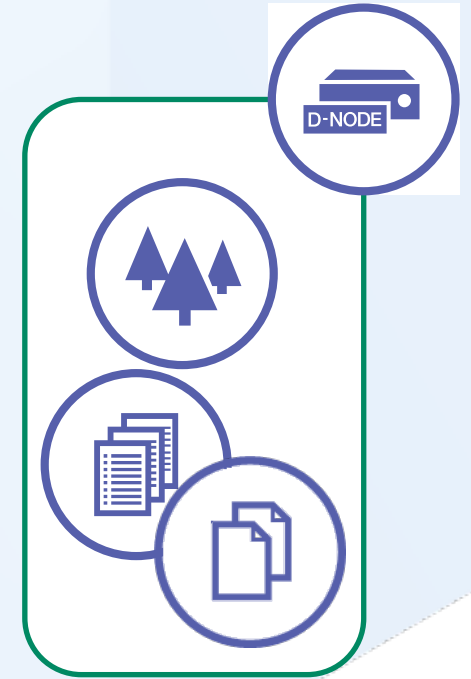
- Exposed via an App Server.
- Is a set of configurations for:
  - What information should be tracked (Indexes/Word Lists).
  - How documents are to be distributed to forests (assignment policy).
  - Backup/Restore operations.
  - Connectivity to multiple forests (usually).





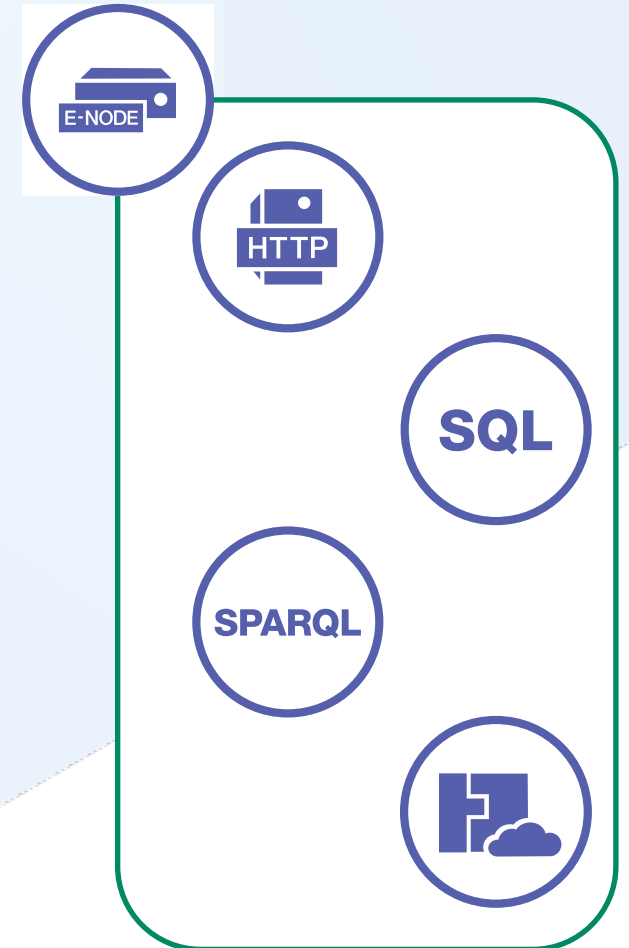
# Forests

- Is the actual unit of storage for documents.
  - A host and location is specified.
  - A forest contains documents and their associated indexes/term lists.
  - A host can have multiple forests.
- A forest can be attached to **one and only one** database.
- A “partition” is a group of forests.
  - Forests in a partition can belong to different hosts.
- More information about forests is available in MarkLogic [docs](#).



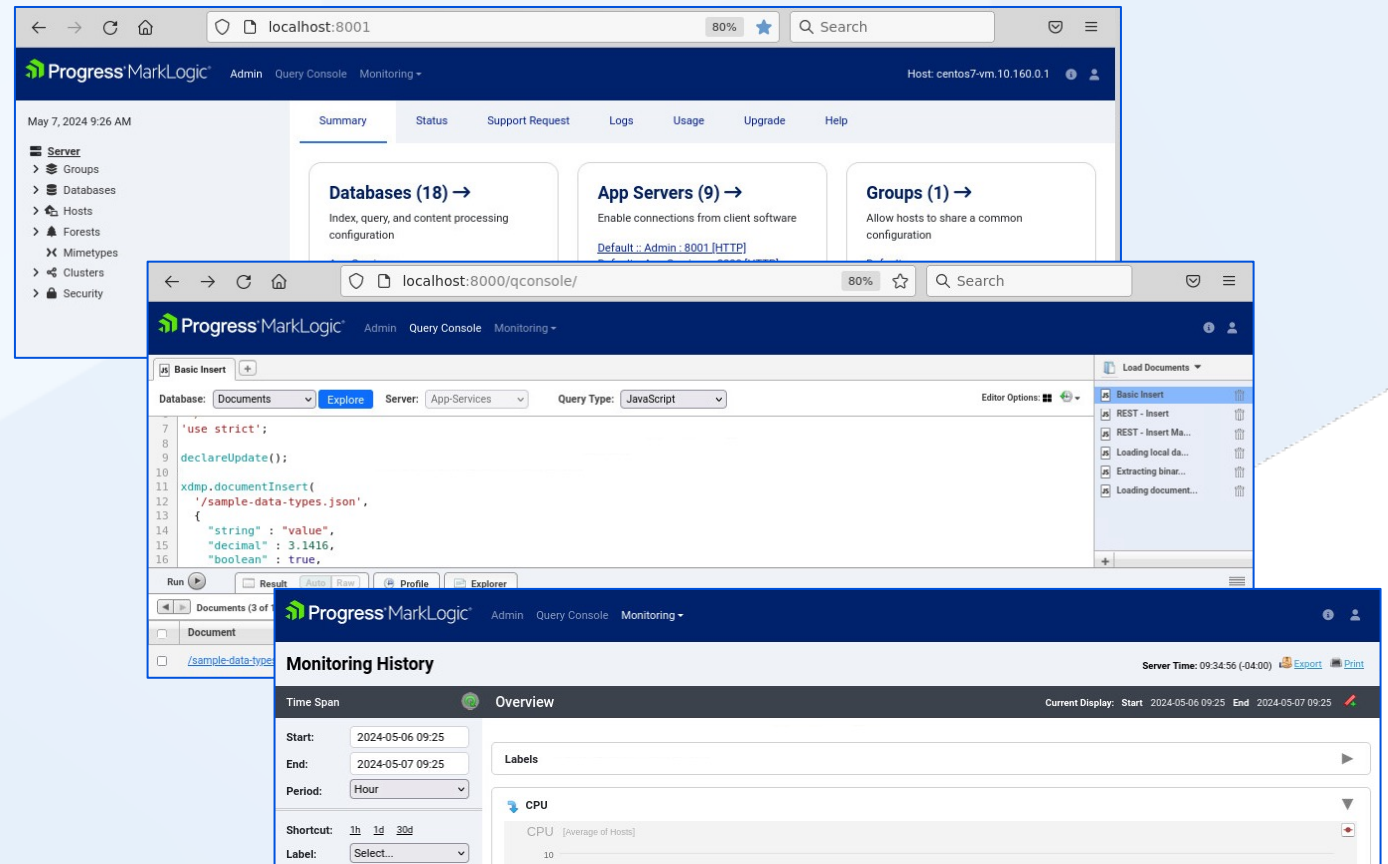
# App Servers

- There are different types of [App Servers](#) for different purposes:
  - HTTP: Access modules via HTTP protocol, e.g. Admin UI (8001)
    - REST: A subtype of HTTP that makes use of special “url rewriters”.
  - ODBC: Meant for your BI tools and ODBC connectors.
  - XDBC: Typically used by MLCP and CORB2.
  - WebDav: Use MarkLogic like a network drive.



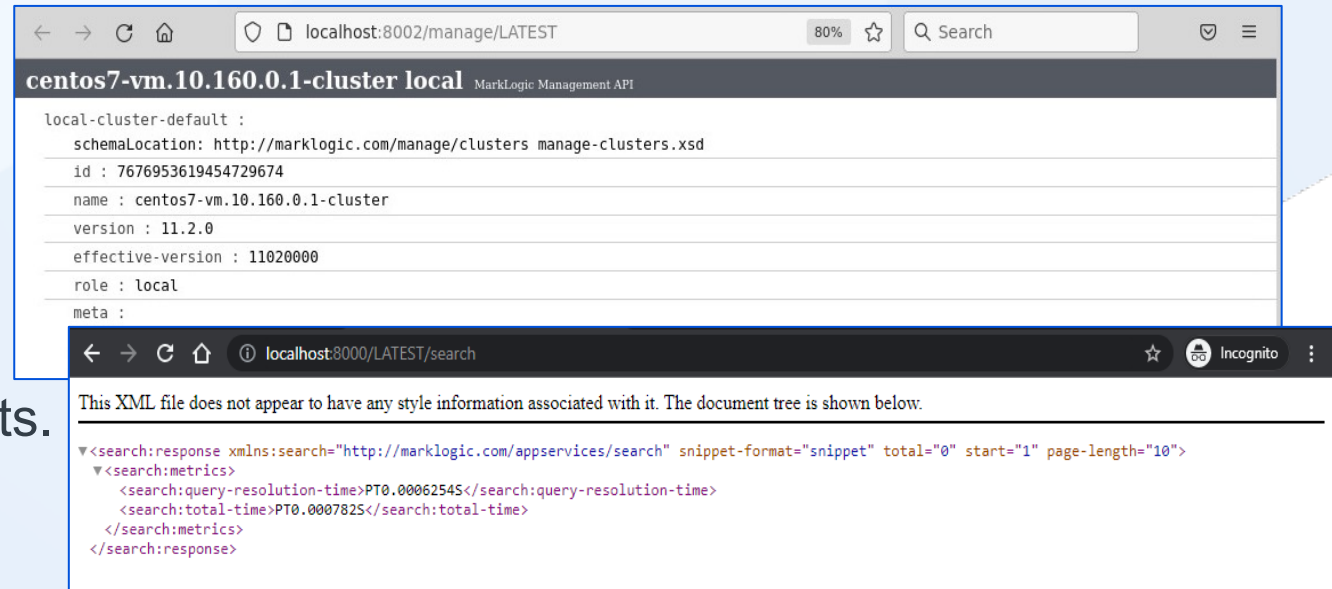
# HTTP App Server

- Accessed using HTTP.
- Runs modules from a specified location or a modules database.
- Built in App Servers:
  - Admin UI (8001)
  - Query Console (8000)
  - Monitoring (8002)



# REST API App Server

- Specialized HTTP App Server:
  - Uses special rewriters and error handlers:
    - `/MarkLogic/rest-api/rewriter.xml`
    - `/MarkLogic/rest-api/error-handler.xqy`
- Supports all [Client API](#) endpoints.
- Built in REST API App Servers:
  - App-Services (8000/LATEST)
  - Manage (8002/manage/LATEST)
    - Supports [Management API](#) endpoints.



# ODBC App Server

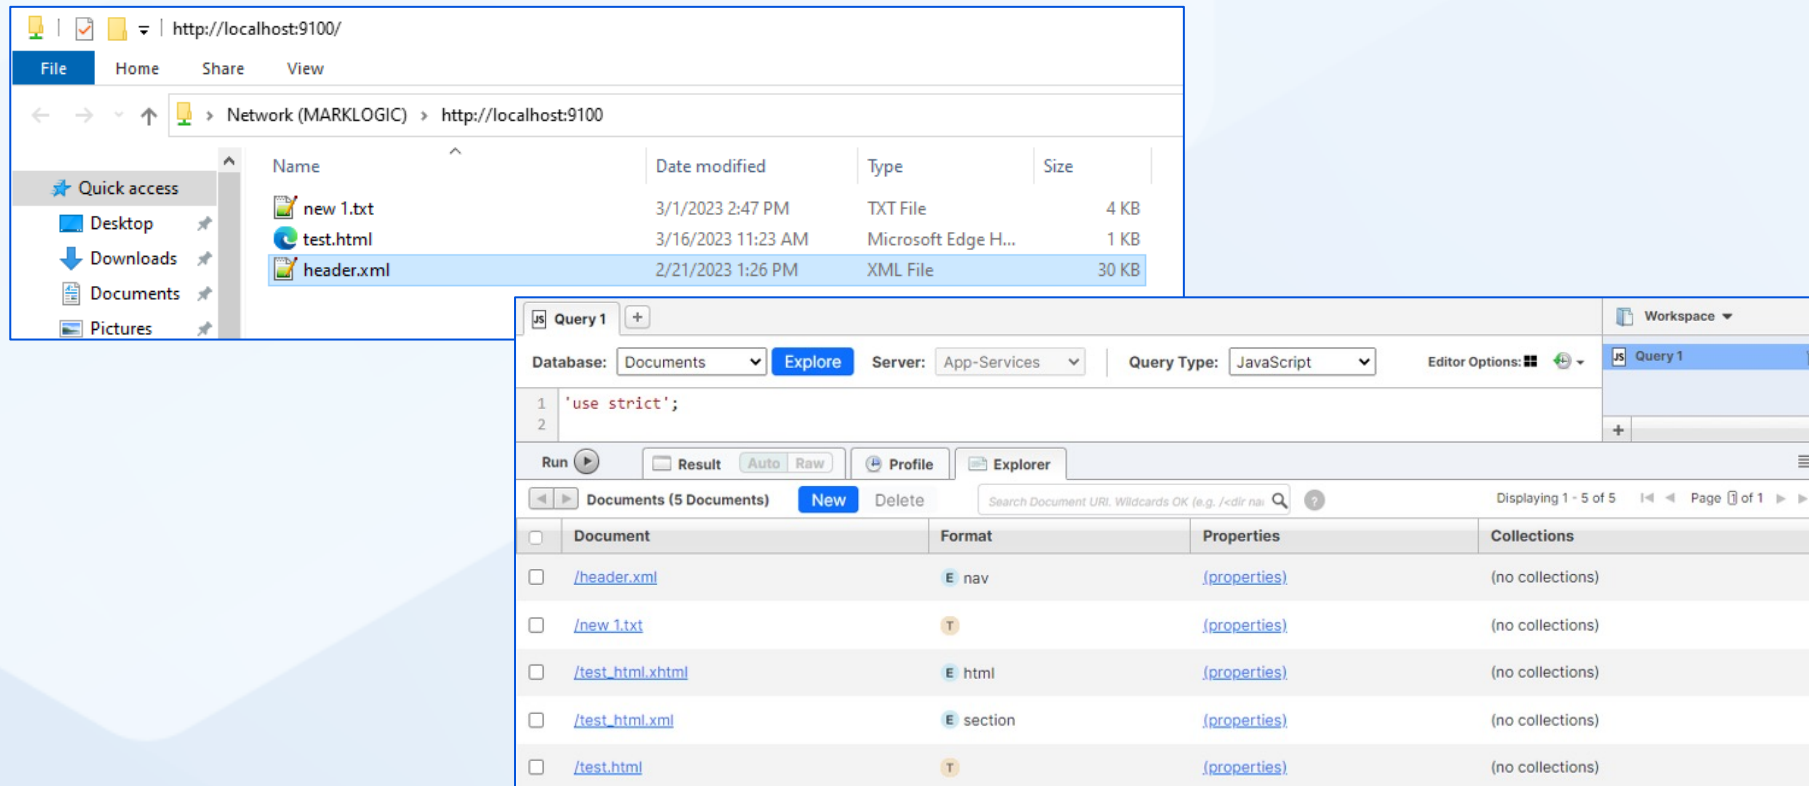
- Supports the connection of Business Intelligence (BI) tools to the MarkLogic Server.
  - Connectors for Power BI and Tableau.
- Can return relational-style data in response to SQL Queries.

# XDBC App Server

- Allows XML Contentbase Connector (XCC) applications to communicate with the MarkLogic Server.
- Used by CORB2, MLCP and similar apps.

# WebDav App Server

- Read and write access a MarkLogic database from a folder.
- Load documents into a MarkLogic database by dragging and dropping into a local or network folder.



The image shows two overlapping windows. The top window is a Windows File Explorer showing a WebDAV folder at `http://localhost:9100/`. The folder contains three files: `new.1.txt` (4 KB), `test.html` (1 KB), and `header.xml` (30 KB). The bottom window is the MarkLogic Explorer interface. It shows a query editor with the text `'use strict';` and a document list table.

Document	Format	Properties	Collections
<a href="#">/header.xml</a>	E nav	<a href="#">(properties)</a>	(no collections)
<a href="#">/new.1.txt</a>	T	<a href="#">(properties)</a>	(no collections)
<a href="#">/test.html.xhtml</a>	E html	<a href="#">(properties)</a>	(no collections)
<a href="#">/test.html.xml</a>	E section	<a href="#">(properties)</a>	(no collections)
<a href="#">/test.html</a>	T	<a href="#">(properties)</a>	(no collections)



**Labs:**

**Create a Forest and Database**

**Create an App Server**





# Configuration Deployment

# Configuration Deployment Challenges

- In the previous labs the Admin UI was used for configuring:
  - a single app server, database and forest.
- Challenge:
  - Correct and Error free Deployment of:
    - This setup across 100 hosts in a single cluster.
    - Indexes, replication, sql views, across MarkLogic DEV, QA, UAT and Production clusters.
- Deployment Tool:
  - ml-gradle (recommended)

# Deployment Tool: ml-gradle

- ml-gradle is:
  - MarkLogic's Gradle plug-in that can quickly “stub out” a new project.
    - **`gradle mlNewProject`**
  - A tool to automate deployment across multiple environments.
  - A single command used to deploy all changes:
    - **`gradle mlDeploy`**

```
project
├── build.gradle
├── gradle-dev.properties
├── gradle-local.properties
├── gradle-prod.properties
├── gradle-qa.properties
├── gradle.properties
└── src
    ├── main
    │   ├── ml-config
    │   │   └── rest-api.json
    │   └── databases
    │       └── content-database.json
    └── ...
        └── ml-modules
```

# ml-gradle

- Supports the use of multiple environment configuration files:
  - argument to switch: **`-PenvironmentName=name`**
- **`gradle.properties`** is the base **`Properties`** file.
- Use the **`gradle-<name>.properties`** to add or override properties for a specific target environment (Dev, QA, Prod).

```
project
  gradle-dev.properties
  gradle-local.properties
  gradle-prod.properties
  gradle-qa.properties
  gradle.properties
```

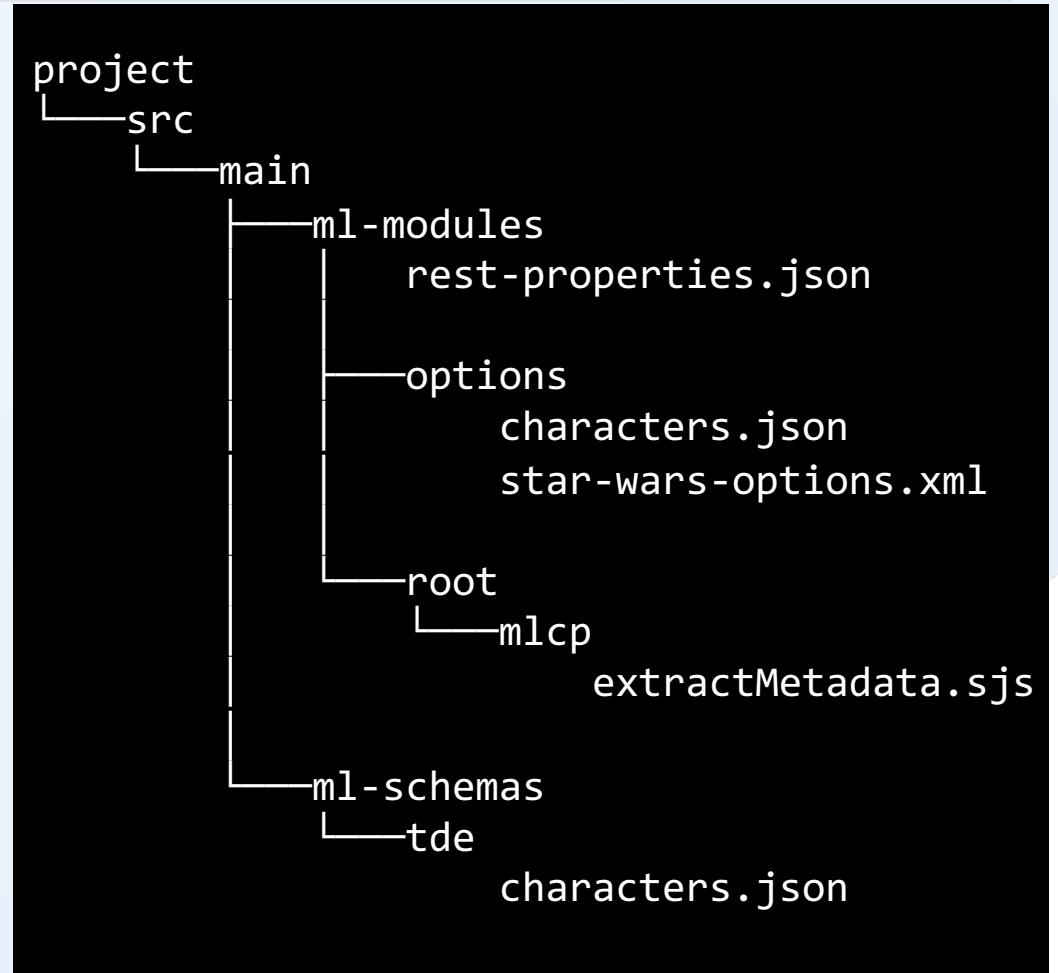
# ml-gradle

- Encode project configuration (databases, indexes, app server, etc.) for version control.
  - This avoids as much “human error” as possible.
- Uses placeholders, `%%propertyName%%`, in these files to capture initialized values from `gradle.properties` during deployment.

```
project
├── src
│   └── main
│       ├── ml-config
│       │   ├── rest-api.json
│       │   ├── databases
│       │   │   ├── content-database.json
│       │   │   └── schema-database.json
│       │   └── servers
│       │       └── odbc.json
```

# ml-gradle

- Used to develop your modules, configure your views with Template Driven Extraction (TDE) and deploy them across all environments.
- Less need to remember which Management REST API to use and how.
- More information about the project layout is available in the [github wiki](#).





## Lab: Use ml-gradle

# Recap

- A host in a cluster with MarkLogic Server installed performs E-node and/or D-node tasks.
- A “database”:
  - Is just a set of configuration parameters.
  - Can have many “forests”.
  - A “forest” can be attached to only one “database”.
- Documents are stored in “forests”.
- An “app server” communicates with a “database”.
- ml-gradle is the recommended deployment tool for MarkLogic implementations.